

Overview

In this assignment you will implement user threads using the Linux pthread API. This will require you to generate new execution contexts (threads) on `pthread_create` and to write a scheduler to switch between your various threads. Since you will be providing a multi-thread environment, you will also need to implement pthread mutexes, mutual exclusion devices that keep a thread locked if it is waiting for a particular mutex. This assignment is intended to illustrate the mechanics and difficulties of scheduling tasks within an operating system.

You will need to do a fair amount of investigation and reading in order to accomplish this assignment.

If you are unfamiliar or uncomfortable with threading and synchronization in C, it will take longer.

This assignment is a group project. Your group may consist of up to 3 students. You should sign up for a group member as soon as possible.

Specifics of Operation

Build a library named "my_pthread_t.h" that contains implementations of the prototypes below.

Pthread Note: Your internal implementation of pthreads should have a running and waiting queue.

Pthreads that are waiting for a mutex should be moved to the waiting queue. Threads that can be scheduled to run should be in the running queue.

```
int my_pthread_create( my_pthread_t * thread, pthread_attr_t * attr, void *(*function)(void*), void * arg);
```

Creates a pthread that executes function. Attributes are ignored, arg is not.

```
void my_thread_yield();
```

Explicit call to the my_thread_t scheduler requesting that the current context can be swapped out and

another can be scheduled if one is waiting.

```
void pthread_exit(void *value_ptr);
```

Explicit call to the my_thread_t library to end the pthread that called it. If the value_ptr isn't NULL,

any return value from the thread will be saved.

```
int my_thread_join(my_thread_t thread, void **value_ptr);
```

Call to the my_thread_t library ensuring that the calling thread will not continue execution until the one it references exits. If value_ptr is not null, the return value of the exiting thread will be passed back.

Mutex note: Both the unlock and lock functions should be very fast. If there are any threads that are meant to compete for these functions, my_thread_yield should be called immediately after running the function in question. Relying on the internal timing will make the function run slower than using yield.

```
int my_thread_mutex_init(my_thread_mutex_t *mutex, const pthread_mutexattr_t *mutexattr);
```

Initializes a my_thread_mutex_t created by the calling thread. Attributes are ignored.

```
int my_thread_mutex_lock(my_thread_mutex_t *mutex);
```

Locks a given mutex, other threads attempting to access this mutex will not run until it is unlocked.

```
int my_pthread_mutex_unlock(my_pthread_mutex_t *mutex);
```

Unlocks a given mutex.

```
int my_pthread_mutex_destroy(my_pthread_mutex_t *mutex);
```

Destroys a given mutex. Mutex should be unlocked before doing so.

It is strongly suggested that you investigate the system library `ucontext.h`. It has a series of commands to make, swap and get the currently running context. When a context is running, it will continue running until it completes. In order to interrupt the current context you should set an interrupt time (setitimer) in 25ms quanta. You should implement and register a signal handler to run any time the interrupt fires. In your signal handler you should schedule and swap to the next context. You can freely modify context and runtime data in your signal handler since, if you are running the signal handler, the signal must have fired, pausing the current context.

Your scheduler should implement a multilevel priority queue, granting highest priority to new threads.

When a thread runs for its entire time slice without terminating, its priority should be decreased. As

threads decrease in priority they should run less often, but for longer. One of the things your scheduler

should get the current system time and check if the current context has run for its entire time slice. If it

has run its entire time slice, it should be swapped out and its priority decreased. If it has not yet run its

entire time slice, no change should be made and the current context should resume computing. If a

thread explicitly yields, a new thread should be scheduled without decreasing the priority of the previous thread. The number of queues, priority levels and how much to increase and decrease priority levels is up to you. Whatever design you pick should be documented in your readme, along with your rationale.

You should establish a maintenance cycle. After some amount of time, when your scheduler runs, you

should scan through all threads in your queue and determine how long they have been running. The

oldest threads should have their priorities scaled up based on how long they have been running. This is

so that the oldest threads do not get starved out of run time. The number of different priority values,

how often you update them, how long the time slices are, how large the time slices are per priority

value, and other details are entirely yours to specify. You should write a suite of benchmarking

functions to test your scheduler with to determine how well your scheduler works. The details of your

scheduler and your testing results should be written up in a readme.pdf

Submission and Structure

Your code should take the form of a `my_pthread.h` header and associated `my_pthread.c` library. Your header file should have a series of macros that replaces all calls to `pthread` and `mutex` function with your own functions. A simple way to delineate them is to prepend “`my_`” to the function names, e.g. `my_pthread_create`. You should support the same API as `pthread`. In particular, your functions should take the same number and type of parameters as `pthread`, with the exception of `pthread` itself. For instance, `my_pthread_create` should take the address of a `my_pthread_t`, not a `pthread_t`.

Compress together your project files into one Asst1.tgz and submit that. Be sure to include your header and library; my_pthread.h and my_pthread.c, as well as a readme.pdf that explains your code's design and operation as well as your scheduling parameters.

Note the names and NetIDs of the students in your group and the ilab machine you wrote and tested your code on as comment at the top of your readme file.

Submit a tarred, gzipped file named Asst1.tgz that contains at least:

my_pthread_t.h

my_pthread.c

readme.pdf

Resources

A POSIX thread library tutorial:

<https://computing.llnl.gov/tutorials/pthreads/>

Another POSIX thread library tutorial:

<http://www.mathcs.emory.edu/~cheung/Courses/455/Syllabus/5c-pthreads/pthreads-tut2.html>

Some notes on implementing thread libraries in Linux:

<http://www.evanjones.ca/software/threading.html>

Please refer to the attached code.tar.gz for a skeleton makefile, library file and header